

Practical-6: Develop a Gradient descent of linear regression using sample dataset.



Output:

```
⇒ Initial cost (theta = zeros): 792.2487672102134
```

```
⇒ Iteration 1: Cost = 680.5572602312128  
Iteration 2: Cost = 585.2192664846582  
Iteration 3: Cost = 503.84035572724207  
Iteration 4: Cost = 434.3766827006044  
Iteration 5: Cost = 375.08365557468636  
Iteration 6: Cost = 324.4721201970103  
Iteration 7: Cost = 281.2709597070975  
Iteration 8: Cost = 244.3951701983308  
Iteration 9: Cost = 212.91861064152818  
Iteration 10: Cost = 186.05074267945125
```

```
⇒ Optimized parameters: theta_0 = 0.08569249990135377, theta_1 = 0.9988176396752504
```

```
⇒ Final cost after optimization: 29.364303963905186
```



Code:

```
import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

# 1. Load dataset

df= pd.read_csv("Experience-Salary.csv")

# Step 2: Extract input (x) and output (y)

data x = data['exp(in months)'].values
y = data['salary(in thousands)'].values
m = len(x) # Number of data points

# Step 3: Prepare the data (add column of ones for intercept term)

X = np.c_[np.ones(m), x] # Shape: (m, 2)

# Step 4: Initialize parameters (theta_0 and theta_1) theta = np.zeros(2) # [theta_0, theta_1]

# Step 5: Learning rate and number of iterations for gradient descent

learning_rate = 0.0001

iterations = 1000

# Step 6: Define the cost function (Mean Squared Error)

def cost_function(X, y, theta):

    predictions = X.dot(theta)

    cost = (1 / m) * np.sum((predictions - y) ** 2)

    return cost

# Step 7: Gradient descent function to optimize theta
```

```
def gradient_descent(X, y, theta, learning_rate, iterations):  
    cost_history = []  
  
    for i in range(iterations):  
        predictions = X.dot(theta)  
        errors = predictions - y  
        gradient = (1 / m) * X.T.dot(errors)  
        theta -= learning_rate * gradient  
        current_cost = cost_function(X, y, theta)  
        cost_history.append(current_cost)  
  
        # Optional: print cost at each iteration  
        print(f"Iteration {i+1}: Cost = {current_cost}")  
  
    return theta, cost_history  
  
# Step 8: Print initial cost before training  
  
initial_cost = cost_function(X, y, theta)  
print(f"Initial cost (theta = zeros): {initial_cost}")  
  
  
# Step 9: Run gradient descent  
  
theta_optimized, cost_history = gradient_descent(X, y, theta, learning_rate, iterations)  
  
# Step 10: Output optimized parameters  
  
print(f"Optimized parameters: theta_0 = {theta_optimized[0]}, theta_1 = {theta_optimized[1]}")  
  
# Step 11: Print final cost  
  
final_cost = cost_function(X, y, theta_optimized)  
print(f"Final cost after optimization: {final_cost}")
```

Step 12: Plot cost function history

```
plt.plot(range(iterations), cost_history)
plt.xlabel('Iterations')
plt.ylabel('Cost (J)')
plt.title('Cost Function History during Gradient Descent') plt.show()
```

Step 13: Plot data points and regression line

```
plt.scatter(x, y, color='red', label='Data points')
plt.plot(x, X.dot(theta_optimized), label='Regression Line', color='blue')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.title('Linear Regression: Experience vs Salary')
plt.legend()
plt.show()
```